

# 操作系统lab5：RV64 缺页异常处理与fork机制

教师：李环、柳晴

学号：3220104519

姓名：蔡佳伟

## 一、实验目的

- 通过 `vm_area_struct` 数据结构实现对进程多区域虚拟内存的管理
- 在 Lab4 实现用户态程序的基础上，添加缺页异常处理 `page fault handler`
- 为进程加入 `fork` 机制，能够支持通过 `fork` 创建新的用户态进程

## 二、实验过程

### 2.1 准备工程

从仓库同步 `user/main.c` 文件并删除原来的 `getpid.c`

修改 `user/Makefile`：

```
user > M Makefile
1  ASM_SRC    = $(filter-out uapp.S, $(sort $(wildcard *.S)))
2  C_SRC      = $(sort $(wildcard *.c))
3  OBJ        = $(patsubst %.S,%.o,$(ASM_SRC)) $(patsubst %.c,%.o,$(C_SRC))
4  TEST       = PFH1
5  CFLAG      = -march=$(ISA) -mabi=$(ABI) -mmodel=medany -fno-builtin -ffunction-sections -fdata-sections
6
```

补充SYS\_CLONE的宏，不然会一直显示报错的，虽然在fork阶段才用到

```
arch > riscv > include > C syscall.h > ...
1  #ifndef __SYSCALL_H__
2  #define __SYSCALL_H__
3  #include "stdint.h"
4  #include "stddef.h"
5  #include "defs.h"
6
7  #define SYS_WRITE 64
8  #define SYS_GETPID 172
9  #define SYS_CLONE 220
10
11 int sys_write(unsigned int fd, const char *buf, size_t count);
12 int sys_getpid();
13 uint64_t do_fork(struct pt_regs *regs);
14 #endif
```

### 2.2 缺页异常处理

## 2.2.1 实现虚拟内存管理功能

先添加vma的定义与数据结构

```
arch > riscv > include > C defs.h > VM_READ
```

```
1  #ifndef __DEFS_H__
2  #define __DEFS_H__
3
4  #include "stdint.h"
5
6  // lab5 vma
7  #define VM_ANON 0x1
8  #define VM_READ 0x2
9  #define VM_WRITE 0x4
10 #define VM_EXEC 0x8
11
```

```
arch > riscv > include > C proc.h > ...
```

```
1  #ifndef PROC_H
15 #define PRIORIT_MAX 10
16
17 #define task_struct_ptr struct task_struct *
18
19 // lab 5
20 struct vm_area_struct
21 {
22     struct mm_struct *vm_mm;           // 所属的 mm_struct
23     uint64_t vm_start;                 // VMA 对应的用户态虚拟地址的开始
24     uint64_t vm_end;                 // VMA 对应的用户态虚拟地址的结束
25     struct vm_area_struct *vm_next, *vm_prev; // 链表指针
26     uint64_t vm_flags;                 // VMA 对应的 flags
27     // struct file *vm_file;           // 对应的文件 (目前还没实现, 而且我们只有一个 uapp 所以暂不需要)
28     uint64_t vm_pgoff; // 如果对应了一个文件, 那么这块 VMA 起始地址对应的文件内容相对文件起始位置的偏移量
29     uint64_t vm_filesz; // 对应的文件内容的长度
30 };
31
32 struct mm_struct
33 {
34     struct vm_area_struct *mmap;
35 };
```

```
37 /* 线程状态段数据结构 */
38 struct thread_struct
39 {
40     uint64_t ra; // 32
41     uint64_t sp; // 40
42     uint64_t s[12]; // 48
43     uint64_t sepc, sstatus, sscratch; // 144 152 160
44 };
45
46 /* 线程数据结构 */
47 struct task_struct
48 {
49     uint64_t state; // 线程状态 0
50     uint64_t counter; // 运行剩余时间 8
51     uint64_t priority; // 运行优先级 1 最低 10 最高, 16
52     uint64_t pid; // 线程 id 24
53
54     struct thread_struct thread;
55     uint64_t *pgd; // 用户态页表 168
56     struct mm_struct mm; // 176
57 };
58
```

每一个 `vm_area_struct` 都对应于 task 地址空间的唯一连续区间。

为了支持 demand paging, 我们需要支持对 `vm_area_struct` 的添加和查找:

`find_vma` 函数: 实现对 `vm_area_struct` 的查找

- 根据传入的地址 `addr`, 遍历链表 `mm` 包含的 VMA 链表, 找到该地址所在的 `vm_area_struct`
- 如果链表中所有的 `vm_area_struct` 都不包含该地址, 则返回 `NULL`

```
// lab 5 function
struct vm_area_struct *find_vma(struct mm_struct *mm, uint64_t addr)
{
    // 遍历查找
    struct vm_area_struct *vma = mm->mmap;
    while (vma)
    {
        if (vma->vm_start <= addr && addr < vma->vm_end)
            return vma;
        vma = vma->vm_next;
    }
    return NULL;
}
```

`do_mmap` 函数: 实现 `vm_area_struct` 的添加

- 新建 `vm_area_struct` 结构体, 根据传入的参数对结构体赋值, 并添加到 `mm` 指向的 VMA 链表中

```
345 uint64_t do_mmap(struct mm_struct *mm, uint64_t addr, uint64_t len, uint64_t vm_pgoff, uint64_t vm_filesz, uint64_t flags)
346 {
347     // 插入到链表头
348     struct vm_area_struct *mmap = (struct vm_area_struct *)kalloc();
349     mmap->vm_mm = mm;
350     mmap->vm_start = addr;
351     mmap->vm_end = addr + len;
352     mmap->vm_flags = flags;
353     mmap->vm_pgoff = vm_pgoff;
354     mmap->vm_filesz = vm_filesz;
355     mmap->vm_prev = NULL;
356     if (mm->mmap != NULL)
357         mm->mmap->vm_prev = mmap;
358     mmap->vm_next = mm->mmap;
359     mm->mmap = mmap;
360     return addr;
361 }
362
```

## 2.2.2 修改task\_init

本次注释掉之前实验对用户栈, 代码load segment的映射操作 (alloc和create\_mapping).

调用 `do_mmap` 函数, 建立用户 task 的虚拟地址空间信息, 在本次实验中仅包括两个区域:

- 代码和数据区域: 该区域从 ELF 给出的 Segment 起始用户态虚拟地址 `phdr->p_vaddr` 开始, 对应文件中偏移量为 `phdr->p_offset` 开始的部分
- 用户栈: 范围为 `[USER_END - PGSIZE, USER_END)`, 权限为 `VM_READ | VM_WRITE`, 并且是匿名的区域 (`VM_ANON`)

```
void load_program(struct task_struct *task)
{
    INFO("\ntask init - load program");
    Elf64_Ehdr *ehdr = (Elf64_Ehdr *)_sramdisk;
    Elf64_Phdr *phdrs = (Elf64_Phdr *)(_sramdisk + ehdr->e_phoff);

    // uint64_t uapp_size = _eramdisk - _sramdisk;
    // int page_num = ((uapp_size + PGSIZE + 1) / PGSIZE);
    // DEBUG("uapp page size: %llx page_num: %llx", uapp_size, page_num);
    // char *new_uapp = (char *)alloc_pages(page_num);
    printk("enter load_program, e_phnum = %d\n", ehdr->e_phnum);
    for (int i = 0; i < ehdr->e_phnum; ++i)
```

```

{
    Elf64_Phdr *phdr = phdrs + i;
    printk("i = %d, type = %d\n", i, phdr->p_type);
    if (phdr->p_type == PT_LOAD)
    {
        // alloc space and copy content
        uint64_t offset = PGOFFSET(phdr->p_vaddr);
        uint64_t size = offset + phdr->p_memsz;
        uint64_t perm = 0x0; // 初始化
        perm = (phdr->p_flags & PF_X) ? perm | PTE_X : perm;
        perm = (phdr->p_flags & PF_W) ? perm | PTE_W : perm;
        perm = (phdr->p_flags & PF_R) ? perm | PTE_R : perm;
        DEBUG("perm: %llx", perm);
        // char *va = (char *)alloc_pages((size + PGSIZE - 1) / PGSIZE);
        DEBUG("size : %llx", phdr->p_memsz + offset);
        DEBUG("page number : %llx , offset number : %llx", (size + PGSIZE -
1) / PGSIZE, offset);
        // segment 使用的内存就是 [p_vaddr, p_vaddr + p_memsz) 这一连续区间，然后
        将 segment 的内容从 ELF 文件中读入到这一内存区间
        // 为了实现对齐，实际上 0 - offset之间的为无效内容
        // for (int i = 0; i < phdr->p_filesz; i++)
        // va[offset + i] = ((char *)ehdr)[phdr->p_offset + i];
        // DEBUG("va : %llx", va);
        // 将 [p_vaddr + p_filesz, p_vaddr + p_memsz) 对应的 *物理区间* 清零
        // memset(va + offset + phdr->p_filesz, 0, phdr->p_memsz - phdr-
>p_filesz);
        // DEBUG("va: %llx, pa: %llx, size: %llx, perm: %llx", va,
        (uint64_t)va - PA2VA_OFFSET, phdr->p_memsz + offset, perm | PTE_U | PTE_V);
        // create_mapping(task->pgd, phdr->p_vaddr, (uint64_t)va -
        PA2VA_OFFSET, size, perm | PTE_U | PTE_V);

        // lab5
        do_mmap(&task->mm, phdr->p_vaddr, phdr->p_memsz, phdr->p_offset,
        phdr->p_filesz, VM_READ | VM_WRITE | VM_EXEC);
    }
}

// DEBUG("user_start: %llx", USER_START);
// uint64_t user_stack = (uint64_t)alloc_page();
// DEBUG("user_stack: %llx", user_stack);
// create_mapping(task->pgd, USER_END - PGSIZE, user_stack - PA2VA_OFFSET,
        PGSIZE, PTE_U | PTE_V | PTE_R | PTE_W);
        do_mmap(&task->mm, USER_END - PGSIZE, PGSIZE, 0, 0, VM_ANON | VM_READ |
        VM_WRITE);
        task->thread.sepc = ehdr->e_entry;
}

```

task\_init函数也需要修改，修改如下：

```

153 void task_init()
154 {
155     INFO("task_init\n");
156     srand(2024);
157     // 1. 调用 kalloc() 为 idle 分配一个物理页
158     // 2. 设置 state 为 TASK_RUNNING;
159     // 3. 由于 idle 不参与调度, 可以将其 counter / priority 设置为 0
160     // 4. 设置 idle 的 pid 为 0
161     // 5. 将 current 和 task[0] 指向 idle
162     idle = (struct task_struct *)kalloc();
163     if (idle == NULL)
164         printk(RED "Error : kalloc() failed in task_init()\n" CLEAR);
165
166     idle->state = TASK_RUNNING;
167     idle->counter = 0;
168     idle->priority = 0;
169     idle->pid = 0;
170
171     current = idle;
172     task[0] = idle;
173
174     // 1. 参考 idle 的设置, 为 task[1] ~ task[NR_TASKS - 1] 进行初始化
175     // 2. 其中每个线程的 state 为 TASK_RUNNING, 此外, counter 和 priority 进行如下赋值:
176     //     - counter = 0;
177     //     - priority = rand() 产生的随机数 (控制范围在 [PRIORITY_MIN, PRIORITY_MAX] 之间)
178     // 3. 为 task[1] ~ task[NR_TASKS - 1] 设置 thread_struct 中的 ra 和 sp
179     //     - ra 设置为 _dummy (见 4.2.2) 的地址
180     //     - sp 设置为该线程申请的物理页的高地址
181
182     printk("_sramdisk: %llx, _eramdisk: %llx, _sbss: %llx\n", (uint64_t)_sramdisk, (uint64_t)_eramdisk, (uint64_t)_sbss);
183     for (int i = 1; i < nr_tasks; i++)
184     {
185         task[i] = (struct task_struct *)kalloc();
186         task[i]->state = TASK_RUNNING;
187         task[i]->counter = 0;
188         task[i]->priority = PRIORITY_MIN + rand() % (PRIORITY_MAX - PRIORITY_MIN + 1);
189         task[i]->pid = i;
190         task[i]->thread.ra = (uint64_t)_dummy;

```

```

191         task[i]->thread.sp = (uint64_t)task[i] + PGSIZE; // 内核态的栈
192
193         // task[i]->thread.sepc = (USER_START); // in loadprogram
194         task[i]->pgd = (uint64_t *)alloc_page();
195         // 将内核页表 swapper_pg_dir 复制到每个进程的页表中
196         for (int j = 0; j < 512; ++j)
197             task[i]->pgd[j] = swapper_pg_dir[j];
198
199         // set_task_pgd(task[i]);
200         load_program(task[i]);
201         task[i]->thread.sstatus = ((1 << 18) | (1 << 5)); // SUM SPIE
202         task[i]->thread.sscratch = (USER_END);
203
204         printk("task[%d] pid = %d priority = %d\n", i, task[i]->pid, task[i]->priority);
205     }
206
207     printk("...task_init done!\n");
208 }
209

```

## 2.3 实现page fault handler

修改 `trap.c`, 为 `trap_handler` 添加捕获 page fault 的逻辑, 分别需要捕获 12, 13, 15 号异常。

```

19 // 中断和异常处理函数
20 void trap_handler(uint64_t scause, uint64_t sepc, struct pt_regs *regs)
21 {
22     uint64_t stval = regs->stval; // 错误发生时的值
23     if (scause & 0x8000000000000000)
24     {
25         // 处理中断
26         uint64_t interrupt_code = scause & 0x7FFFFFFFFFFFFFFF;
27
28         switch (interrupt_code)
29         {
30             case 5:
31                 clock_set_next_event(); // 设置下一次时钟中断
32                 do_timer(); // 处理定时器中断
33                 break;
34             default:
35                 Err("Unknown interrupt\n"); // 未知中断
36                 break;
37         }
38     }
39     else
40     {
41         // 处理异常
42         uint64_t exception_code = scause & 0x7FFFFFFFFFFFFFFF;
43         switch (exception_code)
44         {
45             case 8:
46                 // 处理用户态系统调用
47                 int sys_call_num = regs->x[17]; // 获取系统调用编号
48                 switch (sys_call_num)
49                 {
50                     case SYS_WRITE:
51                         sys_write((unsigned int)regs->x[10], (const char *)regs->x[11], (size_t)regs->x[12]); // 处理write系统调用
52                         break;
53                     case SYS_GETPID:
54                         regs->x[10] = sys_getpid(); // 处理getpid系统调用
55                         break;
56
57                     case SYS_CLONE:
58                         regs->x[10] = do_fork(regs); // 处理fork系统调用
59                         break;
60                     default:
61                         break;
62                 }
63                 regs->sepc += 4; // 更新程序计数器，跳过系统调用指令
64                 break;
65             case 12: // 处理加载页面错误
66                 do_page_fault(regs, 12);
67                 break;
68             case 13: // 处理读取页面错误
69                 do_page_fault(regs, 13);
70                 break;
71             case 15: // 处理写入页面错误
72                 do_page_fault(regs, 15);
73                 break;
74             default:
75                 break;
76         }
77     }
78 }

```

实现缺页异常的函数 `do_page_fault`，具体逻辑如下：

1. 通过 `stval` 获得访问出错的虚拟内存地址（Bad Address）
2. 通过 `find_vma()` 查找 bad address 是否在某个vma 中
  - 如果不在，则出现非预期错误，可以通过 `Err` 宏输出错误信息
  - 如果在，则根据vma 的 flags 权限判断当前page fault 是否合法
    - 如果非法（比如触发的是 instruction page fault 但 vma 权限不允许执行），则 `Err` 输出错误信息
    - 其他情况合法，需要我们按接下来的流程创建映射
3. 分配一个页，接下来要将这个页映射到对应的用户地址空间
4. 通过 `(vma->vm_flags & VM_ANON)` 获得当前的 VMA 是否是匿名空间
  - 如果是匿名空间，则直接映射即可
  - 如果不是，则需要根据 `vma->vm_pgoff` 等信息从 ELF 中读取数据，填充后映射到用户空间

```

79 // 处理页错误 (页异常)
80 void do_page_fault(struct pt_regs *regs, uint64_t exception_code)
81 {
82     Log("[PID = %d PC = %lx] valid page fault at `%lx` with cause %d", current->pid, regs->sepc, regs->stval, exception_code);
83
84     uint64_t bad_addr = regs->stval; // 错误地址
85     if (current->mm.mmap == NULL)
86         Err("current task's mmap is NULL"); // 如果当前任务没有内存映射, 报错
87
88     struct vm_area_struct *vma = find_vma(&current->mm, bad_addr); // 查找错误地址所在的虚拟内存区域 (VMA)
89     if (vma == NULL)
90         Err("No vma found for va: %lx with exception: %d, sepc: %lx\n", bad_addr, exception_code, regs->sepc); // 如果没有找到相应的VMA, 报错
91
92     // 如果访问的内存区域没有读权限, 则报错
93     if (vma->vm_flags & VM_READ == 0)
94         Err("vma perm err without VM_READ: va: %lx\n", bad_addr);
95
96     // 根据异常代码处理不同的情况
97     switch (exception_code)
98     {
99     case 12: // 执行错误
100         if (vma->vm_flags & VM_EXEC == 0) // 如果VMA没有执行权限, 报错
101             Err("vma perm err without VM_EXEC: va: %lx\n", bad_addr);
102         break;
103     case 13: // 加载错误 (读取内存)
104         break;
105     case 15: // 存储错误 (写入内存)
106         if (vma->vm_flags & VM_WRITE == 0) // 如果VMA没有写权限, 报错
107             Err("vma perm err without VM_WRITE: va: %lx\n", bad_addr);
108         break;
109     }
110
111     // 创建新的内存页映射
112     uint64_t *page = (uint64_t *)alloc_page(); // 分配一页新的内存
113     memset((void *)page, 0, PGSIZE); // 将新分配的内存页清零
114

```

```

115     uint64_t perm = (vma->vm_flags & 0b01110) | 0b10001; // 设置内存权限
116
117     if (vma->vm_flags & VM_ANON) // 如果是匿名映射, 直接创建映射
118         create_mapping(current->pgd, PGROUNDNDOWN(bad_addr), (uint64_t)page - PA2VA_OFFSET, PGSIZE, perm);
119     else
120     {
121         // 如果是文件映射, 填充数据
122         uint64_t mem_size = vma->vm_end - vma->vm_start; // 计算内存区域大小
123
124         uint64_t global_part_start = vma->vm_start + vma->vm_filesz; // 文件数据起始位置
125         uint64_t global_part_end = vma->vm_end; // 文件数据结束位置
126         uint64_t bad_seg_start = vma->vm_pgoff + (uint64_t) sramdisk; // 文件段起始位置
127         uint64_t bad_seg_now = bad_seg_start + PGROUNDNDOWN(bad_addr) - vma->vm_start; // 当前映射段的起始位置
128         uint64_t size = PGSIZE; // 要映射的大小
129         uint64_t offset; // 页内的偏移
130
131         if (bad_addr > vma->vm_start + vma->vm_filesz)
132         {
133             offset = 0;
134             size = PGSIZE;
135         }
136
137         // 如果地址小于文件映射区域的起始地址, 需要处理并拷贝数据
138         if (PGROUNDUP(vma->vm_start) > bad_addr)
139         {
140             offset = PGOFFSET(bad_seg_start); // 计算偏移量
141             size = PGSIZE;
142             memcpy((void *) (page), (void *) (bad_seg_now), size); // 将文件数据拷贝到新分配的页中
143         }
144         // 处理其他边界情况
145         else if (PGROUNDNDOWN(vma->vm_start + vma->vm_filesz) == PGROUNDNDOWN(bad_addr))
146         {
147             offset = 0;
148             memcpy((void *) (page + offset / 8), (void *) (bad_seg_now), PGOFFSET(vma->vm_start + vma->vm_filesz));
149         }
150         else if (PGROUNDNDOWN(vma->vm_end) == PGROUNDNDOWN(bad_addr))

```

```

152         size = PGOFFSET(vma->vm_end);
153     }
154
155     // 映射内存页面
156     create_mapping(current->pgd, PGROUNDNDOWN(bad_addr), (uint64_t)page - PA2VA_OFFSET, size, PTE_U | PTE_V | perm);
157 }
158

```

## 2.4 测试缺页处理

使用 `make run TEST=PFH1` 在 `task_init` 中并未进行映射, 直到 `page fault` 触发用户态进程的拷贝和映射, 且只有第一次触发

```

[INFO] setup vm starting
[INFO] setup vm done
...buddy init done!
...mm init done!
[INFO] setup vm final starting
[vm.c,125,create_mapping] root : fffffffe00020c000, [80200000, 80204000] -> [ffffffe000200000, fffffffe000204000], perm: b
[vm.c,125,create_mapping] root : fffffffe00020c000, [80204000, 80206000] -> [ffffffe000204000, fffffffe000206000], perm: 3
[vm.c,125,create_mapping] root : fffffffe00020c000, [80206000, 88000000] -> [ffffffe000206000, fffffffe008000000], perm: 7
[DEBUG] satp_val: 8000000000000000
[INFO] setup vm final done
[INFO] task init

_sramdisk: fffffffe000207000, _eramdisk: fffffffe000208c50, _sbss: fffffffe000209000
[INFO]
task init - load program
enter load_program, e_phnum = 3
i = 0, type = 1879048195
i = 1, type = 1
[DEBUG] perm: e
[DEBUG] size : 23f0
[DEBUG] page number : 3 , offset number : e8
i = 2, type = 1685382481
task[1] pid = 1 priority = 7
...task init done!
2024 ZJU Operating System
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at `100e8` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000] -> [10000, 11000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 10178] valid page fault at `3fffffff8` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000] -> [3fffffff000, 4000000000], perm: 17
[trap.c,82,do_page_fault] [PID = 1 PC = 10198] valid page fault at `12000` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d83f0] -> [12000, 123f0], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 110ac] valid page fault at `110ac` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000] -> [11000, 12000], perm: 1f
[U-MODE] pid: 1, sp is 0x3fffffff0, this is print No.1
[U-MODE] pid: 1, sp is 0x3fffffff0, this is print No.2
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7

```

使用 `make run TEST=PFH2` 类似

```

[INFO] setup vm starting
[INFO] setup vm done
...buddy init done!
...mm init done!
[INFO] setup vm final starting
[vm.c,125,create_mapping] root : fffffffe00020c000, [80200000, 80204000] -> [ffffffe000200000, fffffffe000204000], perm: b
[vm.c,125,create_mapping] root : fffffffe00020c000, [80204000, 80206000] -> [ffffffe000204000, fffffffe000206000], perm: 3
[vm.c,125,create_mapping] root : fffffffe00020c000, [80206000, 88000000] -> [ffffffe000206000, fffffffe008000000], perm: 7
[DEBUG] satp_val: 8000000000000000
[INFO] setup vm final done
[INFO] task init

_sramdisk: fffffffe000207000, _eramdisk: fffffffe000208c48, _sbss: fffffffe000209000
[INFO]
task init - load program
enter load_program, e_phnum = 3
i = 0, type = 1879048195
i = 1, type = 1
[DEBUG] perm: e
[DEBUG] size : 33f8
[DEBUG] page number : 4 , offset number : e8
i = 2, type = 1685382481
task[1] pid = 1 priority = 7
...task init done!
2024 ZJU Operating System
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at `100e8` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000] -> [10000, 11000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 10178] valid page fault at `3fffffff8` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000] -> [3fffffff000, 4000000000], perm: 17
[trap.c,82,do_page_fault] [PID = 1 PC = 10194] valid page fault at `13000` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d83f8] -> [13000, 133f8], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 11080] valid page fault at `11080` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000] -> [11000, 12000], perm: 1f
[U-MODE] pid: 1, increment: 0
[U-MODE] pid: 1, increment: 1
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
[U-MODE] pid: 1, increment: 2
[U-MODE] pid: 1, increment: 3

```

## 3.1 实现fork系统调用 准备工作

修改proc.c相关代码，使得其只初始化一个进程，其他进程保留为 NULL 等待 fork 创建

- 定义nr\_tasks记录当前进程数并作为栈顶指针
- task\_init：只初始化一个进程,将nr\_tasks置为2指向栈顶
- Schedule:将NR\_TASKS替换为nr\_tasks

```

for (int i = 1; i < nr_tasks; i++)
{
    task[i] = (struct task_struct *)kalloc();
    task[i]->state = TASK_RUNNING;
    task[i]->counter = 0;
    task[i]->priority = PRIORITY_MIN + rand() % (PRIORITY_MAX - PRIORITY_MIN + 1);
}

```



```
// lab 5 fork
// 记录当前进程数，并作为 tasks 的栈顶指针来使用
uint64_t nr_tasks = 2;
```

添加系统调用处理

在syscall.h加入#define SYS\_CLONE 220(之前已经加入了)

在trap\_handler调用do\_fork，处理regs->a7 == SYS\_CLONE的情况

```
case 8:
    // 处理用户态系统调用
    int sys_call_num = regs->x[17]; // 获取系统调用编号
    switch (sys_call_num)
    {
        case SYS_WRITE:
            sys_write((unsigned int)regs->x[10], (const char *)regs->x[11], (size_t)regs->x[12]); // 处理write系统调用
            break;
        case SYS_GETPID:
            regs->x[10] = sys_getpid(); // 处理getpid系统调用
            break;
        case SYS_CLONE:
            regs->x[10] = do_fork(regs); // 处理fork系统调用
            break;
        default:
            break;
    }
    regs->sepc += 4; // 更新程序计数器，跳过系统调用指令
    break;
```

## 3.2 do\_fork实现

fork的实现：

- 创建一个新进程：
  - 拷贝内核栈（包括了 `task_struct` 等信息）
  - 创建一个新的页表
    - 拷贝内核页表 `swapper_pg_dir`
    - 遍历父进程 vma，并遍历父进程页表
      - 将这个 vma 也添加到新进程的 vma 链表中
      - 如果该 vma 项有对应的页表项存在（说明已经创建了映射），则需要深拷贝一整页的内容并映射到新页表中

```
39 uint64_t do_fork(struct pt_regs *regs)
40 {
41     Log("current: thread sp : %lx, thread sscratch : %lx", current->thread.sp, current->thread.sscratch);
42     Log("sp : %lx", (uint64_t)regs);
43
44     struct task_struct * _task = (struct task_struct *)alloc_page(); // 为新进程分配一页内存
45     uint64_t cur_sscratch = csr_read(sscratch); // 获取当前的sscratch寄存器值
46     Log("cur_sscratch : %lx", cur_sscratch);
47
48     // 复制父进程的内存映射结构
49     memcpy((void *)_task, (void *)current, PGSIZE);
50     _task->pid = nr_tasks; // 设置新进程的PID
51     nr_tasks++; // 增加全局任务计数
52
53     INFO("PID = %d] forked from [PID = %d]", _task->pid, current->pid); // 记录新进程创建信息
54
55     if (_task->pid >= NR_TASKS)
56         Err("task number exceed"); // 如果任务数量超过上限，报错
57
58     task[_task->pid] = _task; // 将新进程添加到任务数组中
59     _task->mm.mmap = NULL; // 清空新进程的内存映射结构
60
61     // 设置新进程的返回地址和栈指针
62     _task->thread.ra = (uint64_t)_ret_from_fork;
63     _task->thread.sp = (uint64_t)_task + PGSIZE - (35 * 8); // 设置新进程的内核栈位置
64     _task->thread.sscratch = cur_sscratch; // 设置新进程的sscratch寄存器
65
66     // 创建新进程的页表
67     _task->pgd = (uint64_t)alloc_page();
68     memcpy((void *)_task->pgd, (void *)swapper_pg_dir, PGSIZE); // 将内核页表拷贝到新进程的页表中
69 }
```

- 将新进程加入调度队列
- 处理父子进程的返回值
  - 父进程通过 `do_fork` 函数直接返回子进程的 pid，并回到自身运行
  - 子进程通过被调度器调度后（跳到 `thread.ra`），开始执行并返回 0

```

69 // 遍历父进程的虚拟内存区域 (VMA)
70 struct vm_area_struct *pvma = current->mm.mmap;
71 for (; pvma != NULL; pvma = pvma->vm_next)
72 {
73     // 根据父进程的VMA信息，映射相应的内存区域到新进程
74     do_mmap(&_task->mm, pvma->vm_start, pvma->vm_end - pvma->vm_start, pvma->vm_pgoff, pvma->vm_filesz, pvma->vm_flags);
75
76     // 遍历VMA中的每一页，复制父进程的物理页到新进程
77     for (uint64_t va = PGROUNDDOWN(pvma->vm_start); va < pvma->vm_end; va += PGSIZE)
78     {
79         uint64_t pte_res = findall_pte(current->pgd, va); // 获取父进程页表项
80         uint64_t pa = PTE2PA(pte_res); // 获取物理地址
81         if (pte_res == 0)
82             continue;
83         uint64_t *page = (uint64_t *)alloc_page(); // 为新页分配内存
84         memcpy((void *)page, (void *)pa + PA2VA_OFFSET, PGSIZE); // 复制父进程的页到新进程
85
86         Log("copy page : %lx -> %lx", (uint64_t)pa, (uint64_t)page);
87         create_mapping(_task->pgd, va, (uint64_t)page - PA2VA_OFFSET, PGSIZE, PTE_U | PTE_V | pvma->vm_flags); // 创建新进程的页表映射
88     }
89 }
90
91 // 更新新进程的寄存器状态，设置正确的返回值和SEPC
92 struct pt_regs *new_regs = (struct pt_regs *)((PGOFFSET((uint64_t)regs) + (uint64_t)_task));
93 memcpy((void *)new_regs, (void *)regs, sizeof(struct pt_regs));
94 new_regs->x[10] = 0; // 设置a0寄存器为0
95 new_regs->sepc = new_regs->sepc + 4; // 更新SEPC，跳过fork指令
96
97 return _task->pid; // 返回新进程的PID
98
99

```

## 3.3 处理进程返回逻辑

修改entry.S如下：

```

55 sd t0, 272(sp)
56
57 # 2. call trap_handler
58 csrr a0, scause
59 csrr a1, sepc
60 addi a2, sp, 0
61 call trap_handler
62
63 # lab5 fork 后的进程返回地址
64 .globl __ret_from_fork
65
66 __ret_from_fork:
67 # 3. restore sepc and 32 registers (x2(sp) should be restore last) from stack
68 ld t0, 272(sp)

```

此部分还有设置do\_fork的父进程、子进程的用户栈和内核栈指针部分，在前面已经展示

## 3.4 测试fork

make run TEST=FORK1

fork时为已映射页拷贝并映射,fork出的子进程与父进程global\_variable相互独立

fork在PID=2 forked from PID=1后调用完cause = 1f的 create\_mapping 后完成

```

...buddy init done!
...mm init done!
[INFO] setup_vm final starting
[vm.c,125,create_mapping] root : ffffffff00020c000, [80200000, 80204000] -> [ffffffe000200000, fffffffe000204000], perm: b
[vm.c,125,create_mapping] root : fffffffe00020c000, [80204000, 80206000] -> [ffffffe000204000, fffffffe000206000], perm: 3
[vm.c,125,create_mapping] root : fffffffe00020c000, [80206000, 88000000] -> [ffffffe000206000, fffffffe008000000], perm: 7
[DEBUG] satp_val: 80000000008020C
[INFO] setup_vm final done
[INFO] task_init

_sramdisk: fffffffe000207000, _eramdisk: fffffffe000208d50, _sbss: fffffffe000209000
[INFO]
task_init - load program
enter load_program, e_phnum = 3
i = 0, type = 1879048195
i = 1, type = 1
[DEBUG] perm: e
[DEBUG] size : 23f0
[DEBUG] page number : 3, offset number : e8
i = 2, type = 1685382481
task[1] pid = 1 priority = 7
...task init done!
2024 ZJU Operating System
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at '100e8' with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000] -> [10000, 11000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101ac] valid page fault at '3fffffff8' with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000] -> [3fffffff000, 4000000000], perm: 17
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffffcc0
[INFO] [PID = 2] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002db000
[vm.c,125,create_mapping] root : fffffffe0002d9000, [802db000, 802dc000] -> [3fffffff000, 4000000000], perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002df000
[vm.c,125,create_mapping] root : fffffffe0002d9000, [802df000, 802e0000] -> [10000, 11000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 1022c] valid page fault at '12000' with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802e2000, 802e23f0] -> [12000, 123f0], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 11118] valid page fault at '11118' with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802e3000, 802e4000] -> [11000, 12000], perm: 1f
[U-PARENT] pid: 1 is running! global variable: 0

```

make run TEST=FORK2

fork出PID=2的子进程实现深拷贝,正确输出父进程改变的字符串及变量值,且后续与父进程独立.

```

...buddy init done!
...mm init done!
[INFO] setup_vm final starting
[vm.c,125,create_mapping] root : fffffffe00020c000, [80200000, 80204000] -> [ffffffe000200000, fffffffe000204000], perm: b
[vm.c,125,create_mapping] root : fffffffe00020c000, [80204000, 80206000] -> [ffffffe000204000, fffffffe000206000], perm: 3
[vm.c,125,create_mapping] root : fffffffe00020c000, [80206000, 88000000] -> [ffffffe000206000, fffffffe008000000], perm: 7
[DEBUG] satp_val: 80000000008020C
[INFO] setup_vm final done
[INFO] task_init

_sramdisk: fffffffe000207000, _eramdisk: fffffffe000208fc0, _sbss: fffffffe000209000
[INFO]
task_init - load program
enter load_program, e_phnum = 3
i = 0, type = 1879048195
i = 1, type = 1
[DEBUG] perm: e
[DEBUG] size : 43f8
[DEBUG] page number : 5, offset number : e8
i = 2, type = 1685382481
task[1] pid = 1 priority = 7
...task init done!
2024 ZJU Operating System
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at '100e8' with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000] -> [10000, 11000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101ac] valid page fault at '3fffffff8' with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000] -> [3fffffff000, 4000000000], perm: 17
[trap.c,82,do_page_fault] [PID = 1 PC = 101d0] valid page fault at '12000' with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d9000] -> [12000, 13000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 112d4] valid page fault at '112d4' with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000] -> [11000, 12000], perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 1043c] valid page fault at '14008' with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802da000, 802da3f0] -> [14000, 143f0], perm: 1f
[U] pid: 1 is running! global variable: 0
[U] pid: 1 is running! global variable: 1
[U] pid: 1 is running! global variable: 2
[trap.c,82,do_page_fault] [PID = 1 PC = 10228] valid page fault at '13008' with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802db000, 802dc000] -> [13000, 14000], perm: 1f

```

```

[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffffcd0
[INFO] [PID = 2] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002df000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802df000, 802e0000] -> [3fffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002e3000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e3000, 802e4000] -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002e6000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e6000, 802e7000] -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002e7000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e7000, 802e8000] -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802db000 -> fffffffe0002e8000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e8000, 802e9000] -> [13000, 14000), perm: 1f
[syscall.c,87,do_fork] copy page : 802da000 -> fffffffe0002e9000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e9000, 802ea000] -> [14000, 15000), perm: 1f
[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[U-PARENT] pid: 1 is running! global_variable: 3
[U-PARENT] pid: 1 is running! global_variable: 4
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 6, priority = 7
switch from 1 to [PID = 2 PRIORITY = 7 COUNTER = 6]
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[U-CHILD] pid: 2 is running! global_variable: 3
[U-CHILD] pid: 2 is running! global_variable: 4
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
pid = 2, counter = 7, priority = 7
switch from 2 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[U-CHILD] pid: 1 is running! global_variable: 5
[U-CHILD] pid: 1 is running! global_variable: 6
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 7, priority = 7
switch from 1 to [PID = 2 PRIORITY = 7 COUNTER = 7]

```

make run TEST=FORK3

进程能够正确实现多层fork(),并且子进程正确复制父进程global\_variable值,后续父子进程相互独立

```

...buddy init done!
...mm init done!
[INFO] setup_vm_final starting
[vm.c,125,create_mapping] root : fffffffe00020c000, [80200000, 80204000] -> [fffffffe000200000, ffffffffe000204000), perm: b
[vm.c,125,create_mapping] root : fffffffe00020c000, [80204000, 80206000] -> [fffffffe000204000, ffffffffe000206000), perm: 3
[vm.c,125,create_mapping] root : fffffffe00020c000, [80206000, 88000000] -> [fffffffe000206000, ffffffffe008000000), perm: 7
[DEBUG] satp_val: 800000000000020C
[INFO] setup_vm_final done
[INFO] task_init

_sramdisk: fffffffe000207000, _eramdisk: fffffffe000208d28, _sbss: fffffffe000209000
[INFO]
task_init - load program
enter load_program, e_phnum = 3
i = 0, type = 1879048195
i = 1, type = 1
[DEBUG] perm: e
[DEBUG] size : 23f0
[DEBUG] page number : 3 , offset number : e8
i = 2, type = 1685382481
task[1] pid = 1 priority = 7
...task_init done!
2024 ZJU Operating System
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at `100e8` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000] -> [10000, 11000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101ac] valid page fault at `3fffffff8` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000] -> [3fffffff000, 4000000000), perm: 17
[trap.c,82,do_page_fault] [PID = 1 PC = 101c8] valid page fault at `12000` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d83f0] -> [12000, 123f0), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 11134] valid page fault at `11134` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000] -> [11000, 12000), perm: 1f
[U] pid: 1 is running! global_variable: 0
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffffcd0
[INFO] [PID = 2] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002dd000
[vm.c,125,create_mapping] root : fffffffe0002db000, [802dd000, 802de000] -> [3fffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002e1000

```

```

[vm.c,125,create_mapping] root : fffffffe0002db000, [802e1000, 802e2000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002e4000
[vm.c,125,create_mapping] root : fffffffe0002db000, [802e4000, 802e5000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002e5000
[vm.c,125,create_mapping] root : fffffffe0002db000, [802e5000, 802e6000) -> [12000, 13000), perm: 1f
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 3] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002e9000
[vm.c,125,create_mapping] root : fffffffe0002e7000, [802e9000, 802ea000) -> [3fffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002ed000
[vm.c,125,create_mapping] root : fffffffe0002e7000, [802ed000, 802ee000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002f0000
[vm.c,125,create_mapping] root : fffffffe0002e7000, [802f0000, 802f1000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002f1000
[vm.c,125,create_mapping] root : fffffffe0002e7000, [802f1000, 802f2000) -> [12000, 13000), perm: 1f
[U] pid: 1 is running! global_variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 4] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002f5000
[vm.c,125,create_mapping] root : fffffffe0002f3000, [802f5000, 802f6000) -> [3fffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002f9000
[vm.c,125,create_mapping] root : fffffffe0002f3000, [802f9000, 802fa000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002fc000
[vm.c,125,create_mapping] root : fffffffe0002f3000, [802fc000, 802fd000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002fd000
[vm.c,125,create_mapping] root : fffffffe0002f3000, [802fd000, 802fe000) -> [12000, 13000), perm: 1f
[U] pid: 1 is running! global_variable: 2
[U] pid: 1 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 6, priority = 7
pid = 3, counter = 6, priority = 7
pid = 4, counter = 6, priority = 7
switch from 1 to [PID = 2 PRIORITY = 7 COUNTER = 6]
[syscall.c,41,do_fork] current: thread sp : fffffffe0002daee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 5] forked from [PID = 2]
[syscall.c,87,do_fork] copy page : 802e1000 -> fffffffe000301000

```

```

[vm.c,125,create_mapping] root : fffffffe0002ff000, [80301000, 80302000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e4000 -> fffffffe000304000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80304000, 80305000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e5000 -> fffffffe000305000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80305000, 80306000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802dd000 -> fffffffe000307000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80307000, 80308000) -> [3fffffff000, 4000000000), perm: 17
[U] pid: 2 is running! global_variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002daee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 6] forked from [PID = 2]
[syscall.c,87,do_fork] copy page : 802e1000 -> fffffffe00030d000
[vm.c,125,create_mapping] root : fffffffe00030b000, [8030d000, 8030e000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e4000 -> fffffffe000310000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80310000, 80311000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e5000 -> fffffffe000311000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80311000, 80312000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802dd000 -> fffffffe000313000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80313000, 80314000) -> [3fffffff000, 4000000000), perm: 17
[U] pid: 2 is running! global_variable: 2
[U] pid: 2 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 6, priority = 7
pid = 4, counter = 6, priority = 7
pid = 5, counter = 6, priority = 7
pid = 6, counter = 6, priority = 7
switch from 2 to [PID = 3 PRIORITY = 7 COUNTER = 6]
[U] pid: 3 is running! global_variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002e6ee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 7] forked from [PID = 3]
[syscall.c,87,do_fork] copy page : 802ed000 -> fffffffe000319000
[vm.c,125,create_mapping] root : fffffffe000317000, [80319000, 8031a000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802f0000 -> fffffffe00031c000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031c000, 8031d000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802f1000 -> fffffffe00031d000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031d000, 8031e000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e9000 -> fffffffe00031f000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031f000, 80320000) -> [3fffffff000, 4000000000), perm: 17

```

## 三、讨论和心得

本次写代码的过程还是比较困难的，战线拉得比较长，细节问题也比较多，很难处理，之前的代码也有很多要注释掉的部分，我对注释掉哪个部分也不算特别清晰，因而踩了一些坑。

还会有莫名其妙的问题，解决起来更是痛苦，只能通过大改代码来解决。



我在这里设置新进程的内核栈位置错误，没有 $-35*8$ ，会在切换进程的时候卡死，因为栈指针的位置不对。子线程跳回的地址在`_traps`的`call trap_handler`之后，所以要手动调整 $-35 * 8$ 。

```
39  uint64_t do_fork(struct pt_regs *regs)
40  {
41      uint64_t cur_sscratch = cur_read(sscratch); // 获取当前的sscratch寄存器值
42      Log("cur_sscratch : %lx", cur_sscratch);
43
44      // 复制父进程的内存映射结构
45      memcpy((void *)_task, (void *)current, PGSIZE);
46      _task->pid = nr_tasks; // 设置新进程的PID
47      nr_tasks++; // 增加全局任务计数
48
49      INFO("[PID = %d] forked from [PID = %d]", _task->pid, current->pid); // 记录新进程创建信息
50
51      if (_task->pid >= NR_TASKS)
52          Err("task number exceed"); // 如果任务数量超过上限，报错
53
54      task[_task->pid] = _task; // 将新进程添加到任务数组中
55      _task->mm.mmap = NULL; // 清空新进程的内存映射结构
56
57      // 设置新进程的返回地址和栈指针
58      _task->thread.ra = (uint64_t)_ret_from_fork;
59      _task->thread.sp = (uint64_t)_task + PGSIZE; // 设置新进程的内核栈位置
60      _task->thread.sscratch = cur_sscratch; // 设置新进程的sscratch寄存器
61  }
```

```
[INFO] [PID = 4] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002f5000
[vm.c,124,create_mapping] root : fffffffe0002f3000, [802f5000, 802f6000) -> [3ffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002f9000
[vm.c,124,create_mapping] root : fffffffe0002f3000, [802f9000, 802fa000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002fc000
[vm.c,124,create_mapping] root : fffffffe0002f3000, [802fc000, 802fd000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002fd000
[vm.c,124,create_mapping] root : fffffffe0002f3000, [802fd000, 802fe000) -> [12000, 13000), perm: 1f
[U] pid: 1 is running! global_variable: 2
[U] pid: 1 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 6, priority = 7
pid = 3, counter = 6, priority = 7
pid = 4, counter = 6, priority = 7
switch from 1 to [PID = 2 PRIORITY = 7 COUNTER = 6]
```

我在fork的最后一开始没有给new regs的`sepc+4`，这样会一直进行fork指令，直到超过`nr_tasks`数目。

```
[U] pid: 4 is running! global_variable: 2
[U] pid: 4 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 0, priority = 7
pid = 4, counter = 0, priority = 7
pid = 5, counter = 6, priority = 7
pid = 6, counter = 6, priority = 7
pid = 7, counter = 6, priority = 7
switch from 4 to [PID = 5 PRIORITY = 7 COUNTER = 6]
[syscall.c,41,do_fork] current: thread sp : fffffffe0002feee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 8] forked from [PID = 5]
[syscall.c,87,do_fork] copy page : 80307000 -> fffffffe000325000
[vm.c,124,create_mapping] root : fffffffe000323000, [80325000, 80326000) -> [3ffffff000, 4000000000), perm: 17
[syscall.c,87,do_fork] copy page : 80301000 -> fffffffe000329000
[vm.c,124,create_mapping] root : fffffffe000323000, [80329000, 8032a000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 80304000 -> fffffffe00032c000
[vm.c,124,create_mapping] root : fffffffe000323000, [8032c000, 8032d000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 80305000 -> fffffffe00032d000
[vm.c,124,create_mapping] root : fffffffe000323000, [8032d000, 8032e000) -> [12000, 13000), perm: 1f
[U] pid: 5 is running! global_variable: 1
[U] pid: 5 is running! global_variable: 2
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 0, priority = 7
pid = 4, counter = 0, priority = 7
pid = 5, counter = 0, priority = 7
pid = 6, counter = 6, priority = 7
pid = 7, counter = 6, priority = 7
pid = 8, counter = 6, priority = 7
switch from 5 to [PID = 6 PRIORITY = 7 COUNTER = 6]
[syscall.c,41,do_fork] current: thread sp : fffffffe00030aee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 9] forked from [PID = 6]
[syscall.c,56,do_fork] task number exceed
```

## 四、思考题

### 1. 呈现出你在 page fault 的时候拷贝 ELF 程序内容的逻辑。

通过vma获取bad\_addr所在segment和page的起始地址

```
126      uint64_t bad_seg_start = vma->vm_pgoff + (uint64_t)_sramdisk; // 文件段起始位置
127      uint64_t bad_seg_now = bad_seg_start + PGROUNDOWN(bad_addr) - vma->vm_start; // 当前映射段的起始位置
```

拷贝起始地址：

- 如果bad\_seg\_now和bad\_addr在同一页，从seg\_start开始拷贝
- bad\_seg\_now处于一个整页page，直接拷贝所在页的内容
- 如果bad\_seg\_now和file size的末尾在同一个page，只对file size范围内部分进行拷贝，但建立映射的范围是整个page
- bad\_seg\_now和vm\_end在同一个page，从vm\_end所在页的起始部分开始拷贝，size=PGOFFSET(vm\_end)

### 2. 回答4.3.5 中的问题：

- 在 do\_fork 中，父进程的内核栈和用户栈指针分别是什么？

内核栈指针为thread\_struct.sp, 用户栈指针为thread\_struct.sscratch

- 在 do\_fork 中，子进程的内核栈和用户栈指针的值应该是什么？

内核栈指针的值为子进程的pt\_regs地址，用户栈指针的值为当前sscratch寄存器的值，即父进程用户态栈指针

- 在 do\_fork 中，子进程的内核栈和用户栈指针分别应该赋值给谁？

用户态进程的thread\_struct.sp和thread\_struct.sscratch

### 3. 为什么要为子进程 pt\_regs 的 sepc 手动加四？

子进程应该和父进程一样回到用户程序的对应位置。父进程在do\_fork返回用户程序经过trap\_handler时，因为发生的是系统调用异常，需要手动+4，避免重复系统调用语句从而死循环。而子进程返回用户程序不经过trap\_handler，为达到相同效果需要在此时为其手动+4，使得\_traps通过sret返回时，回到下一条语句。

### 4. 对于 Fork main #2 (即 FORK2)，在运行时，ZJU OS Lab5 位于内存的什么位置？是否在读取的时候产生了page fault？请给出必要的截图以说明。

位于[p\_vaddr + p\_filesz, p\_vaddr + p\_memsz)

没有发生page fault，在fork进程时父进程已经映射其所在的内存段，子进程也完成复制以及映射

在[12000, 13000)这段里面

```

switch from 0 to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at `100e8` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000) -> [10000, 11000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101ac] valid page fault at `3fffffff8` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000) -> [3fffffff000, 4000000000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101d0] valid page fault at `12000` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d9000) -> [12000, 13000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 112d4] valid page fault at `112d4` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000) -> [11000, 12000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 1043c] valid page fault at `14008` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802da000, 802da3f8) -> [14000, 143f8), perm: 1f
[U] pid: 1 is running! global_variable: 0
[U] pid: 1 is running! global_variable: 1
[U] pid: 1 is running! global_variable: 2
[trap.c,82,do_page_fault] [PID = 1 PC = 10228] valid page fault at `13008` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802db000, 802dc000) -> [13000, 14000), perm: 1f
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffffcc0
[INFO] [PID = 2] forked from [PID = 1]
[syscall.c,87,do_fork] copy page : 802d5000 -> fffffffe0002df000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802df000, 802e0000) -> [3fffffff000, 4000000000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d2000 -> fffffffe0002e3000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e3000, 802e4000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d9000 -> fffffffe0002e6000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e6000, 802e7000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802d8000 -> fffffffe0002e7000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e7000, 802e8000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802db000 -> fffffffe0002e8000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e8000, 802e9000) -> [13000, 14000), perm: 1f
[syscall.c,87,do_fork] copy page : 802da000 -> fffffffe0002e9000
[vm.c,125,create_mapping] root : fffffffe0002dd000, [802e9000, 802ea000) -> [14000, 15000), perm: 1f
[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[U-PARENT] pid: 1 is running! global_variable: 3
[U-PARENT] pid: 1 is running! global_variable: 4
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 6, priority = 7
switch from 1 to [PID = 2 PRIORITY = 7 COUNTER = 6]
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[U-CHILD] pid: 2 is running! global_variable: 3

```

## 5. 画图分析 make run TEST=FORK3 的进程 fork 过程，并呈现出各个进程的 global\_variable 应该从几开始输出，再与你的输出进行对比验证

开始的全局变量值如下：

```

pid 1 - 0
• pid 2 - 1
    - pid 5 - 1
        ° pid 8 - 2
    - pid 6 - 2
• pid 3 - 1
    - pid 7 - 2
• pid 4 - 2

```

和输出的结果相同：



```

[trap.c,82,do_page_fault] [PID = 1 PC = 100e8] valid page fault at `100e8` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d2000, 802d3000) -> [10000, 11000), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 101ac] valid page fault at `3fffffff8` with cause 15
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d5000, 802d6000) -> [3fffffff000, 4000000000), perm: 17
[trap.c,82,do_page_fault] [PID = 1 PC = 101c8] valid page fault at `12000` with cause 13
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d8000, 802d83f0) -> [12000, 123f0), perm: 1f
[trap.c,82,do_page_fault] [PID = 1 PC = 11134] valid page fault at `11134` with cause 12
[vm.c,125,create_mapping] root : fffffffe0002cf000, [802d9000, 802da000) -> [11000, 12000), perm: 1f
[U] pid: 1 is running! global variable: 0
[syscall.c,41,do_fork] current: thread sp : fffffffe0002cf000, thread sscratch : 4000000000
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 2] forked from [PID = 1]

```

```

[INFO] [PID = 5] forked from [PID = 2]
[syscall.c,87,do_fork] copy page : 802e1000 -> fffffffe000301000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80301000, 80302000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e4000 -> fffffffe000304000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80304000, 80305000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e5000 -> fffffffe000305000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80305000, 80306000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802dd000 -> fffffffe000307000
[vm.c,125,create_mapping] root : fffffffe0002ff000, [80307000, 80308000) -> [3fffffff000, 4000000000), perm: 17
[U] pid: 2 is running! global variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002daee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 6] forked from [PID = 2]
[syscall.c,87,do_fork] copy page : 802e1000 -> fffffffe00030d000
[vm.c,125,create_mapping] root : fffffffe00030b000, [8030d000, 8030e000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e4000 -> fffffffe000310000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80310000, 80311000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e5000 -> fffffffe000311000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80311000, 80312000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802dd000 -> fffffffe000313000
[vm.c,125,create_mapping] root : fffffffe00030b000, [80313000, 80314000) -> [3fffffff000, 4000000000), perm: 17

```

```

switch from 2 to [PID = 3 PRIORITY = 7 COUNTER = 6]
[U] pid: 3 is running! global variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002e6ee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 7] forked from [PID = 3]
[syscall.c,87,do_fork] copy page : 802ed000 -> fffffffe000319000
[vm.c,125,create_mapping] root : fffffffe000317000, [80319000, 8031a000) -> [10000, 11000), perm: 1f
[syscall.c,87,do_fork] copy page : 802f0000 -> fffffffe00031c000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031c000, 8031d000) -> [11000, 12000), perm: 1f
[syscall.c,87,do_fork] copy page : 802f1000 -> fffffffe00031d000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031d000, 8031e000) -> [12000, 13000), perm: 1f
[syscall.c,87,do_fork] copy page : 802e9000 -> fffffffe00031f000
[vm.c,125,create_mapping] root : fffffffe000317000, [8031f000, 80320000) -> [3fffffff000, 4000000000), perm: 17

```

```

switch from 3 to [PID = 4 PRIORITY = 7 COUNTER = 6]
[U] pid: 4 is running! global variable: 2
[U] pid: 4 is running! global variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 0, priority = 7
pid = 4, counter = 0, priority = 7
pid = 5, counter = 6, priority = 7
pid = 6, counter = 6, priority = 7
pid = 7, counter = 6, priority = 7
switch from 4 to [PID = 5 PRIORITY = 7 COUNTER = 6]
[U] pid: 5 is running! global variable: 1
[syscall.c,41,do_fork] current: thread sp : fffffffe0002feee8, thread sscratch : 3fffffff0
[syscall.c,42,do_fork] sp : fffffffe0002ceee8
[syscall.c,46,do_fork] cur_sscratch : 3fffffff0
[INFO] [PID = 8] forked from [PID = 5]

```

```
switch from 5 to [PID = 6 PRIORITY = 7 COUNTER = 6]
[U] pid: 6 is running! global_variable: 2
[U] pid: 6 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 0, priority = 7
pid = 4, counter = 0, priority = 7
pid = 5, counter = 0, priority = 7
pid = 6, counter = 0, priority = 7
pid = 7, counter = 6, priority = 7
pid = 8, counter = 6, priority = 7
switch from 6 to [PID = 7 PRIORITY = 7 COUNTER = 6]
[U] pid: 7 is running! global_variable: 2
[U] pid: 7 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 0, priority = 7
pid = 2, counter = 0, priority = 7
pid = 3, counter = 0, priority = 7
pid = 4, counter = 0, priority = 7
pid = 5, counter = 0, priority = 7
pid = 6, counter = 0, priority = 7
pid = 7, counter = 0, priority = 7
pid = 8, counter = 6, priority = 7
switch from 7 to [PID = 8 PRIORITY = 7 COUNTER = 6]
[U] pid: 8 is running! global_variable: 2
[U] pid: 8 is running! global_variable: 3
SCHEDULE(): INFO OF ALL PROCESSES
pid = 0, counter = 0, priority = 0
pid = 1, counter = 7, priority = 7
```